



Faculty Of Computers and Artificial Intelligence

Cairo University

Image processing

Phase (3)

AYA SABRY 2018035

Gehad Mostafa 20180080

Rana Mohamed 20180104

**Submitted to
Eng.Ibrahim Zedan**

May 2022


Abstract

- Face recognition identifies persons on face images or video frames. In a nutshell, a face recognition system extracts features from an input face image and compares them to the features of labeled faces in a database.
- Comparison is based on a **feature similarity** metric and the label of the **most similar database** entry is used to label the input image.
- If the similarity value is below a certain threshold the input image is labeled as **unknown**.
- Comparing two face images to determine if they show the same person is known as face verification.

Explanation of the proposed system

- Detect, transform, and crop faces on input images. This ensures that faces are **aligned** before feeding them into the CNN.
- This preprocessing step is very important for the performance of the neural network.
- Use the CNN to extract **128-dimensional** representations, or *embeddings*, of faces from the aligned input images.
- In embedding space, **Euclidean distance** directly corresponds to a measure of **face similarity**.
- Compare input embedding vectors to labeled embedding vectors in a database. Here, a **support vector machine (SVM) and a KNN classifier**, trained on labeled embedding vectors, play the role of a database.
- The CNN architecture used here is a variant of the inception architecture
- More precisely, it is a variant of the NN4 architecture and identified as nn4.small2 model in the OpenFace project with **128 hidden units** followed by an **L2** normalization layer on top of the convolutional base. (*Siamese Network*)
- These two top layers are referred to as the **embedding layer** from which the 128-dimensional embedding vectors can be obtained.

Siamese Network



$$\mathcal{L} = \max(d(a, p) - d(a, n) + \text{margin}, 0)$$

Triplet Loss Function

- In Siamese networks, we take an input image of a person and find out the encodings of that image, then, we take the same network without performing any updates on weights or biases and input an image of a different person and again predict its encodings.
- We compare these two encodings to check whether there is a **similarity** between the two images.
- These two encodings act as a latent feature representation of the images. Images with the same person have similar features/encodings. Using this, we compare and tell if the two images have the same person or not.
- We train the network by taking an **anchor** image and comparing it with both a **positive** sample and a **negative** sample.
- The **dissimilarity** between the **anchor image and positive** image must be **low** and the **dissimilarity** between the **anchor image and the negative** image must be **high**.

Experimental result

- The **KNN** classifier achieves an accuracy of **96%** on the test set, the **SVM classifier 98%**.
- Let's use the SVM classifier to illustrate face recognition on a single example.
- Classification results should actually be checked whether (a subset of) the database entries of the predicted identity have a distance **less than τ** otherwise one should assign an **unknown** label.

Screenshots from the program running

Epoch 1/10
100/100 [=====] - 74s 579ms/step - loss: 0.7989

Epoch 2/10
100/100 [=====] - 52s 519ms/step - loss: 0.8012

Epoch 3/10
100/100 [=====] - 49s 487ms/step - loss: 0.8011

Epoch 4/10
100/100 [=====] - 45s 449ms/step - loss: 0.8003

Epoch 5/10
100/100 [=====] - 45s 448ms/step - loss: 0.7999


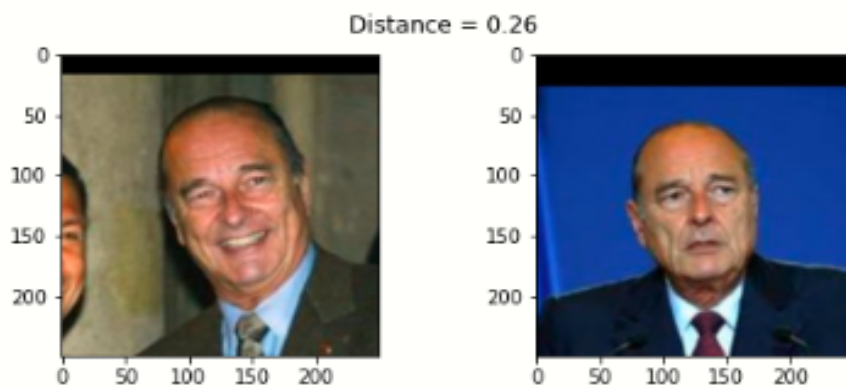
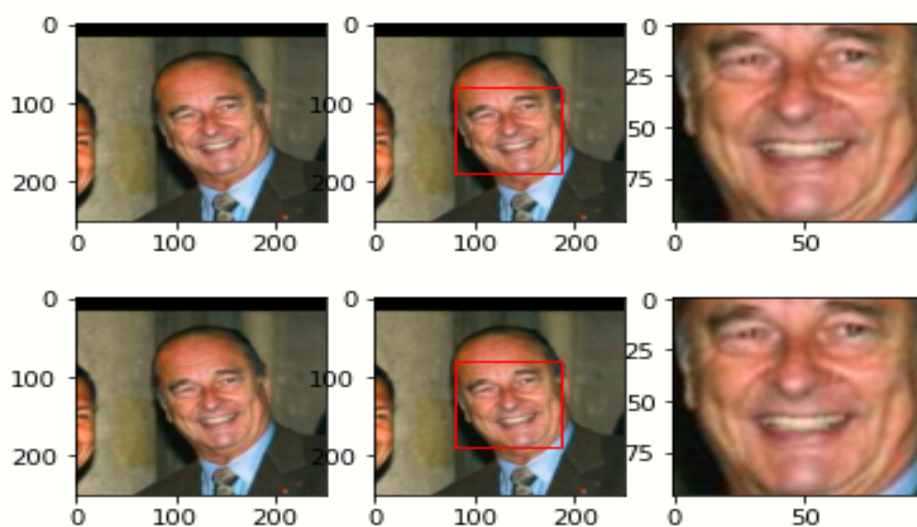
Epoch 6/10
100/100 [=====] - 45s 454ms/step - loss: 0.8004

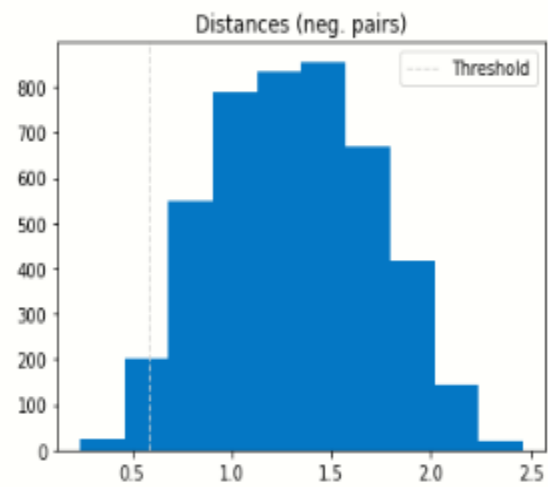


Epoch 7/10
100/100 [=====] - 45s 453ms/step - loss: 0.8000

Epoch 8/10
100/100 [=====] - 45s 454ms/step - loss: 0.8003

Epoch 9/10
100/100 [=====] - 45s 449ms/step - loss: 0.8001

Epoch 10/10
100/100 [=====] - 45s 450ms/step - loss: 0.8001





- KNN accuracy = 0.96, SVM accuracy = 0.98

